

Loan_Approval_Prediction_Model_Using_Neural_Networks (3)

March 23, 2025

```
[132]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from imblearn.over_sampling import SMOTE
from keras import Sequential
from keras.layers import Dense
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
```

```
[133]: from google.colab import auth
auth.authenticate_user()

import gspread
from google.auth import default
creds, _ = default()

gc = gspread.authorize(creds)

worksheet = gc.open('bankloan').sheet1

# get_all_values gives a list of rows.
rows = worksheet.get_all_values()
# print(rows)

# Convert to a DataFrame and render.
import pandas as pd
df = pd.DataFrame.from_records(rows)
```

Data Cleaning

```
[134]: # setting first row as column headers
# dropping old index and resetting new index
# displaying the first 5 rows
df.columns = df.iloc[0]
```

```
df = df.iloc[1:].reset_index(drop=True)
df.head()
```

```
[134]: 0  Loan_ID Gender Married Dependents Education Self_Employed \
0  LP001002 Male No 0 Graduate No
1  LP001003 Male Yes 1 Graduate No
2  LP001005 Male Yes 0 Graduate Yes
3  LP001006 Male Yes 0 Not Graduate No
4  LP001008 Male No 0 Graduate No

0 ApplicantIncome CoapplicantIncome LoanAmount Loan_Amount_Term \
0 5849 0 300 360
1 4583 1508 128 360
2 3000 0 66 360
3 2583 2358 120 360
4 6000 0 141 360

0 Credit_History Property_Area Loan_Status
0 1 Urban Y
1 1 Rural N
2 1 Urban Y
3 1 Urban Y
4 1 Urban Y
```

```
[135]: # checking duplicates
# no duplicates found
df[df.duplicated(keep=False)]
```

```
[135]: Empty DataFrame
Columns: [Loan_ID, Gender, Married, Dependents, Education, Self_Employed,
ApplicantIncome, CoapplicantIncome, LoanAmount, Loan_Amount_Term,
Credit_History, Property_Area, Loan_Status]
Index: []
```

```
[136]: # checking for missing values
# no dupliactes found
df.isnull().sum().sum()
```

```
[136]: np.int64(0)
```

```
[137]: df.isna().sum()
```

```
[137]: 0
Loan_ID 0
Gender 0
Married 0
Dependents 0
```

```

Education      0
Self_Employed  0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount     0
Loan_Amount_Term 0
Credit_History 0
Property_Area   0
Loan_Status     0
dtype: int64

```

```

[138]: # checking all columns that have an empty string
for col in df.columns:
    print(col, df[df[col] == ''].shape[0])

```

```

Loan_ID 0
Gender 0
Married 0
Dependents 0
Education 0
Self_Employed 0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount 0
Loan_Amount_Term 0
Credit_History 0
Property_Area 0
Loan_Status 0

```

```

[139]: df.dtypes

```

```

[139]: 0
Loan_ID      object
Gender       object
Married      object
Dependents   object
Education    object
Self_Employed object
ApplicantIncome object
CoapplicantIncome object
LoanAmount   object
Loan_Amount_Term object
Credit_History object
Property_Area object
Loan_Status  object
dtype: object

```

```
[140]: # changing to the correct data type
df['Dependents'] = df['Dependents'].astype(int)
df['ApplicantIncome'] = df['ApplicantIncome'].astype(float)
df['CoapplicantIncome'] = df['CoapplicantIncome'].astype(float)
df['LoanAmount'] = df['LoanAmount'].astype(float)
df['Loan_Amount_Term'] = df['Loan_Amount_Term'].astype(int)
```

```
[141]: # displaying the shape of the data
df.shape
```

```
[141]: (614, 13)
```

```
[142]: # descriptive statistics for the DataFrame
df.describe()
```

```
[142]: 0      Dependents  ApplicantIncome  CoapplicantIncome  LoanAmount  \
count    614.000000         614.000000         614.000000    614.000000
mean      0.786645         5403.459283         1621.245798    152.257329
std       1.017827         6109.041673         2926.248369     89.356494
min       0.000000         150.000000          0.000000     9.000000
25%       0.000000         2877.500000          0.000000    100.250000
50%       0.000000         3812.500000         1188.500000    129.000000
75%       2.000000         5795.000000         2297.250000    175.000000
max       3.000000        81000.000000        41667.000000   700.000000

0      Loan_Amount_Term
count         614.000000
mean         335.228013
std          78.182120
min          12.000000
25%         360.000000
50%         360.000000
75%         360.000000
max         480.000000
```

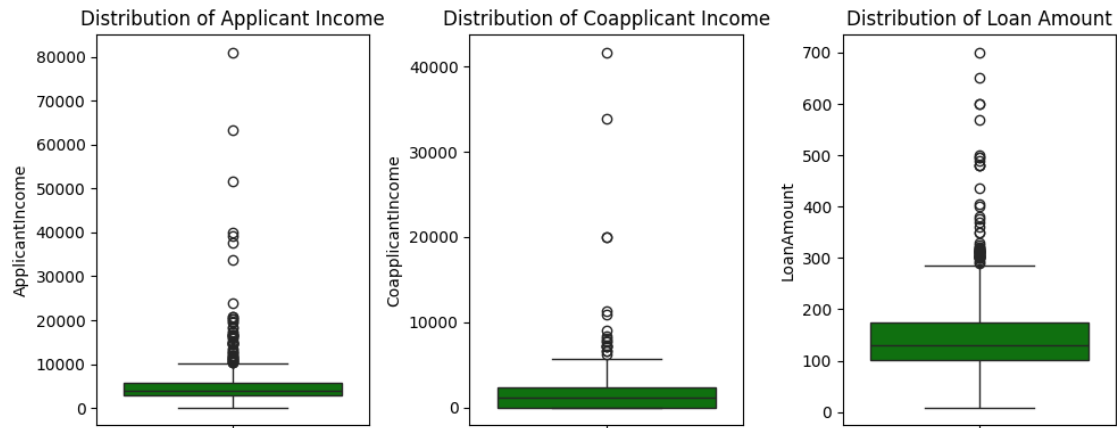
0.1 Exploratory Data Analysis

```
[143]: # box plots showing ditribution of Applicantincome, CoapplicantIncome and Loan_
↳ Amount
fig, ax = plt.subplots(1,3, figsize=(10,4))
sns.boxplot(y = df["ApplicantIncome"], ax= ax[0], color='green')
ax[0].set_title("Distribution of Applicant Income")

sns.boxplot(df["CoapplicantIncome"], ax= ax[1], color='green')
ax[1].set_title("Distribution of Coapplicant Income")

sns.boxplot(df["LoanAmount"], ax= ax[2], color='green')
```

```
ax[2].set_title("Distribution of Loan Amount")
plt.tight_layout()
plt.show()
```



0.1.1 Insights

1. Applicant income is higher and more widely distributed than coapplicant income. This means that the primary financial responsibility for the loan likely rests with the applicant.
2. A significant portion of coapplicants have limited or no income. This could imply the coapplicant's role being more about providing additional security or meeting eligibility criteria rather than contributing financially.
3. Loan amounts are concentrated in a lower range, with fewer instances of very large loans. This may be influenced by the income distribution of the applicants and the lending criteria. This is actually a good sign because it may be in the means of the borrower to pay thus mitigating risks.
4. The presence of outliers in all three distributions highlights the diversity in the applicant pool and the range of loan amounts requested and approved. Instead of removal, they will be analyzed separately.

```
[144]: df_yes = df[df["Loan_Status"] == 'Y']
df_yes.head(2)
```

```
[144]: 0  Loan_ID Gender Married Dependents Education Self_Employed \
0  LP001002  Male      No           0  Graduate             No
2  LP001005  Male     Yes           0  Graduate             Yes

0  ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
0           5849.0             0.0         300.0             360
2           3000.0             0.0         66.0             360

0  Credit_History Property_Area Loan_Status
0              1      Urban           Y
```

2 1 Urban Y

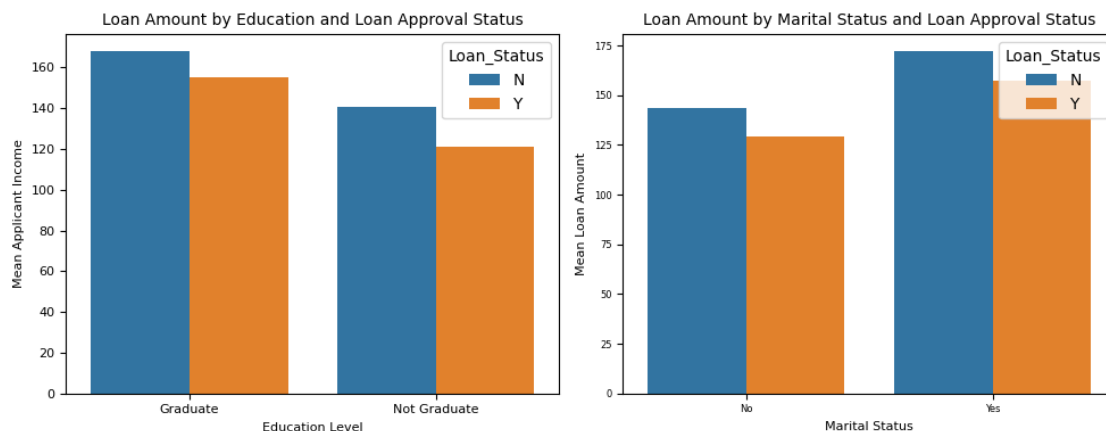
```
[145]: # Group by Education, Loan Status, and Marital Status, and calculate the mean
       ↪ of Applicant Income and Loan Amount
df_educ = df.groupby(['Education', 'Loan_Status'])[['LoanAmount']].mean().
       ↪ reset_index()
df_marital = df.groupby(['Married', 'Loan_Status'])[['LoanAmount']].mean().
       ↪ reset_index()

# Create a subplot with 1 row and 2 columns
fig, ax = plt.subplots(1, 2, figsize=(10, 4))

# Plotting Applicant Income by Education and Loan Status
sns.barplot(data=df_educ, x='Education', y='LoanAmount', hue='Loan_Status',
       ↪ ax=ax[0])
ax[0].set_title('Loan Amount by Education and Loan Approval Status',
       ↪ fontsize=10)
ax[0].set_xlabel('Education Level', fontsize=8)
ax[0].set_ylabel('Mean Applicant Income', fontsize=8)
ax[0].tick_params(axis='both', labelsize=8)

# Plotting Loan Amount by Marital Status and Loan Status
sns.barplot(data=df_marital, x='Married', y='LoanAmount', hue='Loan_Status',
       ↪ ax=ax[1])
ax[1].set_title('Loan Amount by Marital Status and Loan Approval Status',
       ↪ fontsize=10)
ax[1].set_xlabel('Marital Status', fontsize=8)
ax[1].set_ylabel('Mean Loan Amount', fontsize=8)
ax[1].tick_params(axis='both', labelsize=6)

# Adjust layout for better spacing
plt.tight_layout()
plt.show()
```



0.1.2 Insights

Married

1. Disapproved loans: Married individuals whose loan was not approved have a higher mean loan amount compared to unmarried individuals whose loan was not approved.
2. Approved loans: Married individuals whose loan was approved have a higher mean loan amount compared to unmarried individuals whose loan was approved.
3. Overall: Married individuals tend to have a higher mean loan amount compared to unmarried individuals, regardless of the loan approval status. For both marital statuses, the mean loan amount for disapproved loans is higher than for approved loans. The married are more likely to have loan approval

Education Status

1. Education level is positively associated with the mean loan amount. Graduates, on average, apply for and are more likely to be considered for larger loans than non-graduates.

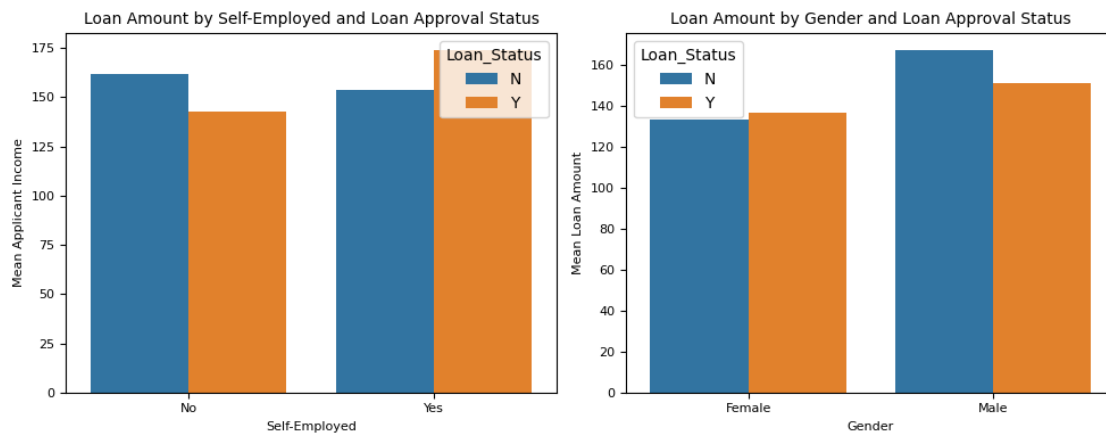
```
[146]: # Group by Self-Employed, Gender, and Loan Status, and calculate the mean of
      ↪ Applicant Income and Loan Amount
df_selfemployed = df.groupby(['Self_Employed', 'Loan_Status'])[['LoanAmount']].
      ↪ mean().reset_index()
df_gender = df.groupby(['Gender', 'Loan_Status'])[['LoanAmount']].mean().
      ↪ reset_index()

# Create a subplot with 1 row and 2 columns for gender and self-employed status
fig, ax = plt.subplots(1, 2, figsize=(10, 4))

# Plotting Applicant Income by Self-Employed and Loan Status
sns.barplot(data=df_selfemployed, x='Self_Employed', y='LoanAmount',
      ↪ hue='Loan_Status', ax=ax[0])
ax[0].set_title('Loan Amount by Self-Employed and Loan Approval Status',
      ↪ fontsize=10)
ax[0].set_xlabel('Self-Employed', fontsize=8)
ax[0].set_ylabel('Mean Applicant Income', fontsize=8)
ax[0].tick_params(axis='both', labelsize=8)

# Plotting Loan Amount by Gender and Loan Status
sns.barplot(data=df_gender, x='Gender', y='LoanAmount', hue='Loan_Status',
      ↪ ax=ax[1])
ax[1].set_title('Loan Amount by Gender and Loan Approval Status', fontsize=10)
ax[1].set_xlabel('Gender', fontsize=8)
ax[1].set_ylabel('Mean Loan Amount', fontsize=8)
ax[1].tick_params(axis='both', labelsize=8)
```

```
plt.tight_layout()
plt.show()
```



0.1.3 Insights

1. Self-employment doesn't show a great impact on the mean loan amount based on approval status. For both self-employed and non-self-employed individuals, the mean loan amount for disapproved loans is similar or slightly higher than for approved loans.
2. For males, a higher mean loan amount is associated with loan disapproval. This could suggest that males might be applying for larger loans that are more frequently rejected, or other factors related to males applying for larger loans lead to higher rejection rates for those amounts.
3. For females, the trend is reversed; a higher mean loan amount is associated with loan approval. This might indicate that when females apply for larger loans, they are more likely to be approved.

0.2 Feature Engineering

```
[147]: # Dropping column Loan ID because its just an indentifier
df.drop('Loan_ID', axis=1, inplace=True)
df.head(2)
```

```
[147]: 0 Gender Married Dependents Education Self_Employed ApplicantIncome \
0 Male No 0 Graduate No 5849.0
1 Male Yes 1 Graduate No 4583.0

0 CoapplicantIncome LoanAmount Loan_Amount_Term Credit_History \
0 0.0 300.0 360 1
1 1508.0 128.0 360 1

0 Property_Area Loan_Status
0 Urban Y
```


1 Rural N

```
[148]: # Identifying independent and dependent variables
```

```
pre_x = df.drop('Loan_Status', axis=1)
```

```
pre_y = df['Loan_Status']
```

```
[149]: # Checking for unique values in y
```

```
pre_y.unique()
```

```
[149]: array(['Y', 'N'], dtype=object)
```

```
[150]: # Convert Loan_Status to binary (1 for 'Y', 0 for 'N')
```

```
dum_y = pre_y.map(dict(Y=1,N=0))
```

```
[151]: # Convert categorical features into dummy variables,
```

```
# dropping the first category to avoid multicollinearity
```

```
x_dum = pd.get_dummies(pre_x, drop_first=True)
```

```
x_dum = x_dum.astype(int)
```

```
x_dum.head(2)
```

```
[151]:
```

	Dependents	ApplicantIncome	CoapplicantIncome	LoanAmount	\
0	0	5849	0	300	
1	1	4583	1508	128	

	Loan_Amount_Term	Gender_Male	Married_Yes	Education_Not Graduate	\
0	360	1	0		0
1	360	1	1		0

	Self_Employed_Yes	Credit_History_1	Property_Area_Semiurban	\
0	0	1		0
1	0	1		0

	Property_Area_Urban
0	1
1	0

```
[152]: # Value counts for dum_y
```

```
# Imbalanced distribution so smote will be applied to balance classes
```

```
dum_y.value_counts()
```

```
[152]: Loan_Status
```

```
1     422
```

```
0     192
```

```
Name: count, dtype: int64
```

```
[153]: smote = SMOTE()
```

```
x_smote, y = smote.fit_resample(x_dum, dum_y)
```

```
min_max = MinMaxScaler()
x = min_max.fit_transform(x_smote)
```

```
[154]: # Counting y values
y.value_counts()
```

```
[154]: Loan_Status
1      422
0      422
Name: count, dtype: int64
```

```
[155]: # Splitting the data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳ random_state = 42, shuffle = True)
```

0.3 Training the Model

```
[156]: # Initialize a Sequential model
classifier = Sequential()
classifier.add(Dense(200, activation='relu', kernel_initializer =
↳ 'random_normal', input_dim = x_test.shape[1])) # Input layer
classifier.add(Dense(400, activation='relu', kernel_initializer =
↳ 'random_normal')) # Hidden layer
classifier.add(Dense(4, activation='relu', kernel_initializer =
↳ 'random_normal'))
classifier.add(Dense(1, activation='sigmoid', kernel_initializer =
↳ 'random_normal')) # Out put layer
classifier.compile(optimizer='adam', loss = 'binary_crossentropy', metrics =
↳ ['accuracy']) # Compiling the model
classifier.fit(x_train, y_train, batch_size = 20, epochs = 50, verbose = 0) #
↳ Training
eval_model = classifier.evaluate(x_train, y_train) # Evaluating the trained
↳ model
eval_model
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
22/22          0s 4ms/step -
accuracy: 0.8391 - loss: 0.3440
```

```
[156]: [0.33888551592826843, 0.8399999737739563]
```

```
[157]: y_predicted = classifier.predict(x_test)
y_predicted = (y_predicted > 0.5).astype(float)
```

6/6 0s 13ms/step

```
[158]: # Convert probabilities into loan status
# Return both credit score (probability) and loan status
loan_status = np.where(y_predicted > 0.5, 'Y', 'N')
predicted_results = list(zip(y_predicted.flatten(), loan_status))
for credit_score, status in predicted_results:
    print(f"Predicted Credit Score: {credit_score:.2f}, Loan Status: {status}")
```

```
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
```

[illegible]

[illegible]

Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']
Predicted Credit Score: 1.00, Loan Status: ['Y']
Predicted Credit Score: 0.00, Loan Status: ['N']

```
[159]: # Accuracy Score
ac = round(accuracy_score(y_test, y_predicted),2)
ac
```


[159]: 0.81

```
[160]: # confusion matrix
cm = confusion_matrix(y_test, y_predicted)
cm
```

```
[160]: array([[67, 27],  
            [ 5, 70]])
```

1. The major diagonal values indicate how well the model is performing in terms of correct classifications. A high sum of these values means the model is making accurate predictions.
2. The model correctly predicted 73 cases as “No” (0).
3. The model incorrectly predicted 21 “No” (0) cases as a “Yes” (1)
4. The model incorrectly predicted 11 cases as “No” (0) when they were actually “Yes” (1).
5. The model correctly predicted 64 cases as “Yes” (1).

```
[161]: # plotting a heat map for confusin matrix  
sns.heatmap(cm, annot=True, fmt='g')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.show()
```



```
[162]: # Pickling the model  
import joblib  
joblib.dump(classifier, 'loan_model.pkl')
```

```
[162]: ['loan_model.pkl']
```

```
[162]:
```