

ML_Scaling_Practice

March 26, 2025

```
[ ]: # Importing the libraries
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder
```

```
[ ]: from google.colab import auth
auth.authenticate_user()

import gspread
from google.auth import default
creds, _ = default()
gc = gspread.authorize(creds)

worksheet = gc.open('cereal-kaggle').sheet1

# get_all_values gives a list of rows.
rows = worksheet.get_all_values()
# print(rows)

# Convert to a DataFrame and render.
import pandas as pd
df = pd.DataFrame.from_records(rows)
```

```
[ ]: # setting first row as headers
# resetting index
# displaying first 5 rows
df.columns = df.iloc[0]
df = df.iloc[1:].reset_index(drop=True)
df.head()
```

```
[ ]: 0          name mfr type calories protein fat sodium fiber carbo \
0          100% Bran   N   C         4     1   130    10     5
1    100% Natural Bran   Q   C        120     3   15     2     8
2          All-Bran    K   C         70     4   1    260     9     7
```

3	All-Bran with Extra Fiber	K	C	50	4	0	140	14	8
4	Almond Delight	R	C		2	2	200	1	14

	sugars	potass	vitamins	shelf	weight	cups	rating
0	6	280	25	top	1	0.33	68.402973
1	8	135	0	top	1	1	33.983679
2	5	320	25	top	1	0.33	59.425505
3	0	330	25	top	1	0.5	93.704912
4	8	-1	25		1	0.75	34.384843

0.0.1 Scale the numeric features

```
[ ]: # mfr, type, calories, protein, fat, fiber, sugars, shelf
df_list = ['mfr', 'type', 'calories', 'protein', 'fat', 'fiber', 'sugars', 'shelf']
df_new = df[df_list]
df_new.head(2)
```

```
[ ]: 0 mfr type calories protein fat fiber sugars shelf
0 N C 120 3 5 2 8 top
1 Q C 120 3 5 2 8 top
```

```
[ ]: # Instantiate a StandardScaler for the numeric features.
scaler = StandardScaler()
```

```
[ ]: # Fit the scaler on X_train_num_imputed
x_num_imputed_list = ['calories', 'protein', 'fat', 'fiber', 'sugars']
x_num_imputed = df_new[x_num_imputed_list]
```

```
[ ]: # Checking the data types
x_num_imputed.dtypes

# Replacing the empty tsrings with np.nan
x_num_imputed = x_num_imputed.replace('', np.nan)

# Changing to float data type
x_num_imputed = x_num_imputed.astype(float)
```

```
[ ]: # Splitting into x_train_num_imputed and x_test_num_imputed
X_train_num_imputed, X_test_num_imputed = train_test_split(x_num_imputed,
    test_size=0.2, random_state=42)
```

```
[ ]: # Fit the scaler on X_train_num_imputed
scaler.fit(X_train_num_imputed)
```

```
[ ]: StandardScaler()
```

```
[ ]: # save the output as "X_train_num_scaled"
X_train_num_scaled = scaler.transform(X_train_num_imputed)
```

```
[ ]: # Transform the numeric test data (X_test_num_imputed) and save the transformed
      ↪data as "X_test_num_scaled"
X_test_num_scaled = scaler.transform(X_test_num_imputed)
```

```
[ ]: # Preview the first 5 rows of X_test_num_scaled.
X_test_num_scaled[:5]
```

```
[ ]: array([[      nan, -0.53311399,  0.95545914, -0.43881613,  0.33525006],
          [      nan, -1.43644603,  0.95545914, -0.43881613,  1.00983861],
          [      nan, -1.43644603,  0.95545914, -0.85709841,  1.23470145],
          [      nan,  1.27355009, -0.01802753,  3.32572435, -0.11447563],
          [      nan,  1.27355009,  1.92894582,  0.39774842,  1.00983861]])
```

```
[ ]: # Describe() of the X_train_num_imputed
X_train_num_imputed.describe()
```

```
[ ]: 0      calories  protein   fat  fiber  sugars
count      61.00    61.00  54.00  55.00   55.00
mean     106.39     2.59   1.02   2.05    6.51
std       20.00     1.12   1.04   2.41    4.49
min       50.00     1.00   0.00   0.00   -1.00
25%      100.00     2.00   0.00   0.00    3.00
50%      110.00     3.00   1.00   1.50    6.00
75%      110.00     3.00   1.75   3.00   10.00
max      160.00     6.00   5.00  14.00   15.00
```

```
[ ]: # Set pandas display option to avoid scientific notation
pd.set_option('display.float_format', '{:.2f}'.format)

# converting X_train_num_scaled into a df
col = X_train_num_imputed.columns.tolist()
X_train_num_scaled_df = pd.DataFrame(X_train_num_scaled, columns=col)
X_train_num_scaled_df.describe()
```

```
[ ]:      calories  protein   fat  fiber  sugars
count      61.00    61.00  54.00  55.00   55.00
mean         0.00    -0.00 -0.00 -0.00    0.00
std          1.01     1.01  1.01  1.01    1.01
min         -2.84    -1.44 -0.99 -0.86   -1.69
25%         -0.32    -0.53 -0.99 -0.86   -0.79
50%          0.18     0.37 -0.02 -0.23   -0.11
75%          0.18     0.37  0.71  0.40    0.78
max          2.70     3.08  3.88  5.00    1.91
```

Did any of the following statistics change after scaling?:Min, Max, Mean, Std Yes, all the above mentioned statistics changed.

This is because when Scaling the mean is set to 0 and standard deviation is set to 1.

During scaling the formular: $x_scaled = (x - \text{mean}) / \text{standard deviation}$ is applied to all the original data points (x).

This transformation alters the distribution of the data, which consequently changes the Min and Max values as well.

0.0.2 Scale the ordinal features

```
[ ]: #ordinal_col = x_train['shelf']  
      #print(ordinal_col)  
      ordinal_col = ['shelf']
```

```
[ ]: df_new[ordinal_col].head()
```

```
[ ]: 0 shelf  
      0 top  
      1 top  
      2 top  
      3 top  
      4
```

```
[ ]: # # One Hot Encoding x_ord  
      x_ord_encoded = pd.get_dummies(df_new[ordinal_col])  
      x_ord_encoded = x_ord_encoded.astype(int)  
      x_ord_encoded.head()
```

```
[ ]:      shelf_  shelf_bottom  shelf_middle  shelf_top  
      0         0             0             0           1  
      1         0             0             0           1  
      2         0             0             0           1  
      3         0             0             0           1  
      4         1             0             0           0
```

```
[ ]: # Splitting x_ord_encoded  
      x_train_ord_encoded, x_test_ord_encoded = train_test_split(x_ord_encoded,  
      ↪test_size=0.2, random_state=42)
```

```
[ ]: # Scaling x_train_ord_encoded  
      ss = StandardScaler()  
      x_train_ord_scaled = scaler.fit_transform(x_train_ord_encoded)
```

```
[ ]: # Transform the ordinal test data (X_test_ord_encoded) and save the transformed  
      ↪data as "X_test_ord_scaled"  
      x_test_ord_scaled = scaler.transform(x_test_ord_encoded)
```

```
[ ]: # Preview the first 5 rows of X_train_ord_encoded.
x_train_ord_encoded.head()
```

```
[ ]:      shelf_  shelf_bottom  shelf_middle  shelf_top
9         0         0         0         1
5         0         1         0         0
34        0         0         0         1
22        0         0         0         1
30        0         1         0         0
```

```
[ ]: # .describe() of X_train_ord_encoded
x_train_ord_encoded.describe()
```

```
[ ]:      shelf_  shelf_bottom  shelf_middle  shelf_top
count    61.00         61.00         61.00         61.00
mean      0.00          0.30          0.23          0.48
std       0.00          0.46          0.42          0.50
min       0.00          0.00          0.00          0.00
25%       0.00          0.00          0.00          0.00
50%       0.00          0.00          0.00          0.00
75%       0.00          1.00          0.00          1.00
max       0.00          1.00          1.00          1.00
```

```
[ ]: # vs. X_train_ord_scaled
x_train_ord_scaled_df = pd.DataFrame(x_train_ord_scaled,
    ↪columns=x_train_ord_encoded.columns)
x_train_ord_scaled_df.describe()
```

```
[ ]:      shelf_  shelf_bottom  shelf_middle  shelf_top
count    61.00         61.00         61.00         61.00
mean      0.00          0.00         -0.00          0.00
std       0.00          1.01          1.01          1.01
min       0.00         -0.65         -0.55         -0.95
25%       0.00         -0.65         -0.55         -0.95
50%       0.00         -0.65         -0.55         -0.95
75%       0.00          1.55         -0.55          1.05
max       0.00          1.55          1.83          1.05
```

Answer the question(s) in a Markdown cell: Did any of the following statistics change after scaling?: Min, Max, Mean, Std? After scaling, the shelf features (which were originally ordinal) are transformed into continuous values with a mean of 0 and a standard deviation of 1, as expected from the StandardScaler. The values are no longer restricted to 0 and 1 but instead spread across a range of negative and positive values.

Cat features

```
[ ]: x_cat = df_new[['mfr', 'type']]
x_cat.head()
```

```
[ ]: 0 mfr type
0    N    C
1    Q    C
2    K    C
3    K    C
4    R    C
```

```
[ ]: x_cat_encoded = pd.get_dummies(x_cat, columns=['mfr', 'type'])
x_cat_encoded = x_cat_encoded.astype(int)
x_cat_encoded.head()
```

```
[ ]:      mfr_  mfr_A  mfr_G  mfr_K  mfr_N  mfr_P  mfr_Q  mfr_R  type_C  type_H
0      0      0      0      0      1      0      0      0      1      0
1      0      0      0      0      0      0      1      0      1      0
2      0      0      0      1      0      0      0      0      1      0
3      0      0      0      1      0      0      0      0      1      0
4      0      0      0      0      0      0      0      1      1      0
```

```
[ ]: # Splitting x_cat_encoded
x_train_cat_encoded, x_test_cat_encoded = train_test_split(x_cat_encoded,
↳ test_size=0.2, random_state=42)
```

Bringing it together

```
[ ]: # Create the final training data, "X_train_processed" by concatenating:
↳ X_train_num_scaled X_train_ord_scaled X_train_cat_encoded
X_train_processed = pd.
↳ concat([X_train_num_scaled_df, x_train_ord_scaled_df, x_train_cat_encoded], axis=1)
X_train_processed.head()
```

```
[ ]:      calories  protein  fat  fiber  sugars  shelf_  shelf_bottom  shelf_middle  \
0      -0.83      0.37  NaN   NaN      NaN      0.00      -0.65      -0.55
1       0.18     -0.53  0.96 -0.23      0.78      0.00       1.55      -0.55
2       0.69      0.37  1.93  0.40     -0.56      0.00      -0.65      -0.55
3     -0.32     -0.53  NaN  -0.02      0.78      0.00      -0.65      -0.55
4     -0.32     -0.53 -0.99 -0.86      1.91      0.00       1.55      -0.55

      shelf_top  mfr_  mfr_A  mfr_G  mfr_K  mfr_N  mfr_P  mfr_Q  mfr_R  type_C  \
0       1.05   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN
1     -0.95  0.00  0.00  0.00  0.00  0.00  0.00  1.00  0.00  1.00
2       1.05  0.00  0.00  0.00  1.00  0.00  0.00  0.00  0.00  1.00
3       1.05  0.00  0.00  0.00  1.00  0.00  0.00  0.00  0.00  1.00
4     -0.95   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN
```

```

type_H
0      NaN
1      0.00
2      0.00
3      0.00
4      NaN

```

```

[ ]: # creating x_test_num_scaled_df
col = X_train_num_imputed.columns.tolist()
X_test_num_scaled_df = pd.DataFrame(X_test_num_scaled, columns=col)

#creating x_test_ord_scaled_df
x_test_ord_scaled_df = pd.DataFrame(x_test_ord_scaled,
    ↪columns=x_train_ord_encoded.columns)

# Create the final test data, "X_test_processed" by concatenating:
    ↪X_test_num_scaled X_test_ord_scaled X_test_cat_encoded
X_test_processed = pd.
    ↪concat([X_test_num_scaled_df,x_test_ord_scaled_df,x_test_cat_encoded],axis=1)
X_test_processed.head()

```

```

[ ]:  calories  protein  fat  fiber  sugars  shelf_  shelf_bottom  shelf_middle \
0      NaN      -0.53  0.96  -0.44   0.34    1.00      -0.65      -0.55
1      NaN     -1.44  0.96  -0.44   1.01    1.00      -0.65      -0.55
2      NaN     -1.44  0.96  -0.86   1.23    0.00      -0.65      1.83
3      NaN     1.27 -0.02   3.33  -0.11    0.00      -0.65      -0.55
4      NaN     1.27  1.93   0.40   1.01    0.00      -0.65      -0.55

```

```

      shelf_top  mfr_  mfr_A  mfr_G  mfr_K  mfr_N  mfr_P  mfr_Q  mfr_R  type_C \
0      -0.95  0.00   0.00   0.00   0.00   1.00   0.00   0.00   0.00   1.00
1      -0.95  NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN
2      -0.95  NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN
3      1.05  NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN    NaN
4      1.05  0.00   0.00   0.00   0.00   0.00   0.00   0.00   1.00   1.00

```

```

type_H
0      0.00
1      NaN
2      NaN
3      NaN
4      0.00

```

```

[ ]: X_test_processed.shape

```

```

[ ]: (28, 19)

```

```
[ ]: X_train_processed.shape
```

```
[ ]: (74, 19)
```

```
[ ]:
```